outlineL08-w4TR-student

**CS 354 - Machine Organization & Programming**
**Tuesday Sept 27 and Thursday Sept 29, 2022**

**Midterm Exam - Thursday, October 6, 7:30 - 9:30 pm**

- **Room: Students will be assigned a room based on their Last "Family" name and Lec**
- **UW ID required**
- **#2 pencils required**
- **closed book, no notes, no electronic devices (e.g., calculators, phones, watches)**
- **see "Midterm Exam 1" on course site Assignments for topics**

**Project p2A:** Due on or before Friday, Sept 30th
**Project p2B:** Due on or before Friday, Oct 7
**Homework hw1:** Due on Monday Sept 26th (solution available Wed morning)
**Homework hw2:** Due on Monday Oct 3rd (solution available Wed morning)

**Last Week**

| | |
|---|---|
| Command-line Arguments<br>Recall 2D Arrays<br>2D Arrays on the Heap<br>2D Arrays on the Stack<br>2D Arrays: Stack vs. Heap<br>Array Caveats | Meet Structures<br>Nesting in Structures and Arrays of Structures |

**This Week**

| | |
|---|---|
| Passing Structures (from L6)<br>Pointers to Structures (from L6)<br>Standard & String I/O in `stdio.h`<br>File I/O in `stdio.h`<br>Copying Text Files<br><br>Three Faces of Memory | Virtual Address Space<br>C's Abstract Memory Model<br><br>Meet Globals and Static Locals<br>Where Do I Live?<br>Linux: Processes and Address Spaces<br>Exam Sample Cover Page |
| **Next Week**: The Heap & Dynamic Memory Allocators (p3)<br>Read: B&O 9.1, 9.2, 9.9.1-9.9.6<br>9.1 Physical and Virtual Addressing<br>9.2 Address Spaces<br>9.9 Dynamic Memory Allocation<br>9.9.1-9.9.6 | |

 **CS 354 (F22): L8 - 1**

## Standard and String I/O in `stdio.h`

### Standard I/O

Standard Input

```
getchar //reads 1 char
gets    //reads 1 string ending with a newline char, BUFFER MIGHT OVERFLOW
```

**`int scanf(const char *format_string, &v1, &v2, ...)`**
reads formatted input from the console keyboard
returns number of inputs stored, or EOF if error/end-of-file occurs before any inputs

*format string*

*format specifiers*

*whitespace*

Standard Output

```
putchar //writes 1 char
puts    //writes 1 string
```

**`int printf(const char *format_string, v1, v2, ...)`**
writes formatted output to the console terminal window
returns number of characters written, or a negative if error

*format string*

Standard Error

**`void perror(const char *str)`**
writes formatted error output to the console terminal window

### String I/O

```
int sscanf(const char *str, const char *format_string, &v1, &v2, ...)
```
reads formatted input from the specified `str`
returns number of characters read, or a negative if error

```
int sprintf(char *str, const char *format_string, v1, v2, ...)
```
writes formatted output to the specified `str`
returns number of characters written, or a negative if error

## File I/O in stdio.h

**Standard I/O Redirection**

**File I/O**

File Input

```
fgetc/getc, ungetc //reads 1 char at a time
fgets              //reads 1 string terminate with a newline char or EOF
int fscanf(FILE *stream, const char *format_string, &v1, &v2, ...)
```

reads formatted input from the specified `stream`
returns number of inputs stored, or EOF if error/end-of-file occurs before any inputs

File Output

```
fputc/putc         //writes 1 char at a time
fputs              //writes 1 string
int fprintf(FILE *stream, const char *format_string, v1, v2, ...)
```

writes formatted output to the specified `stream`
returns number of characters written, or a negative if error

**Predefined File Pointers**

| | |
|---|---|
| `stdin` | is console keyboard |
| `stdout` | is console terminal window |
| `stderr` | is console terminal window, second stream for errors |

**Opening and Closing**

```
FILE *fopen(const char *filename, const char *mode)
```

opens the specified `filename` in the specified `mode`
returns file pointer to the opened file's descriptor, or NULL if there's an access problem

```
int fclose(FILE *stream)
```

flushes the output buffer and then closes the specified `stream`
returns 0, or EOF if error

## Copying Text Files

```
#include <stdio.h>
#include <stdlib.h>

int main(int argc, char *argv[]) {

   if (argc != 3) {
      fprintf(stderr, "Usage: copy inputfile outputfile\n");
      exit(1);
   }


   FILE *ifp =
   if (ifp == NULL) {
      fprintf(stderr, "Can't open input file %s!\n", argv[1]);
      exit(1);
   }


   FILE *ofp =
   if (ofp == NULL) {
      fprintf(stderr, "Can't open output file %s!\n", argv[2]);


      exit(1);
   }


   const int bufsize = 257; //WARNING: assumes lines <= 256 chars
   char buffer[bufsize];









   return 0;
}
```





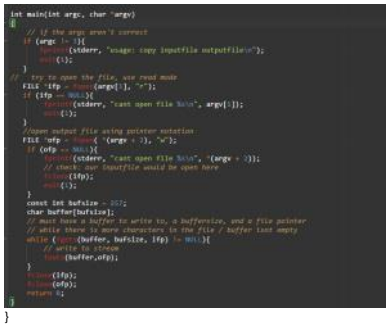

```
int main(int argc, char *argv)
{
   // if the args aren't correct
   if (argc != 3){
      fprintf(stderr, "usage: copy inputfile outputfile\n");
      exit(1);
   }
//   try to open the file, use read mode
   FILE *ifp = fopen(argv[1], "r");
   if (ifp == NULL){
      fprintf(stderr, "cant open file %s\n", argv[1]);
      exit(1);
   }
   //open output file using pointer notation
   FILE *ofp = fopen( *(argv + 2), "w");
   if (ofp == NULL){
      fprintf(stderr, "cant open file %s\n", *(argv + 2));
      // check: our inputfile would be open here
      fclose(ifp);
      exit(1);
   }
   const int bufsize = 257;
   char buffer[bufsize];
   // must have a buffer to write to, a buffersize, and a file pointer
   // while there is more characters in the file / buffer isnt empty
   while (fgets(buffer, bufsize, ifp) != NULL){
      // write to stream
      fputs(buffer,ofp);
   }
   fclose(ifp);
   fclose(ofp);
   return 0;
}
```

## Three Faces of Memory

* *Abstraction:* Manages complexity by hiding unnecessary details

**Process View = Virtual Memory**

Goal: Provide simple view to programmers

*virtual address space (VAS)*: illusion provided by OS: each app has its own memory space

*virtual address*: simulated address that process generates

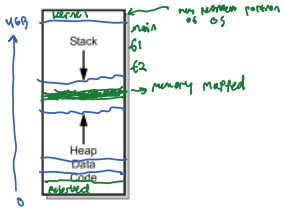


**System View = Illusionist (CS 537)**

Goal: make memory shareable and secure

*pages*: mem blocks managed by OS (4KB)

*page table*: OS data structure that maps virtual pages to physical pages of memory

- secure
- efficient

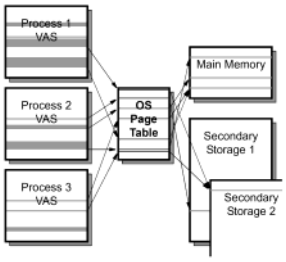


**Hardware View = Physical Memory**

Goal: Keep CPU busy!

*physical address space (PAS)*: multi-level hierarchy, that ensures frequently accessed data is close to CPU

we cover this in 3rd

*physical address*: actual address used to access main chip memory

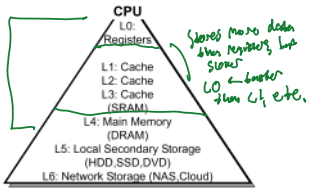

## Virtual Address Space (IA-32/Linux)

**32-bit Processor = 32-bit Addresses => $2^{32}$ = 4,294,967,296 = 4GB Address Space**

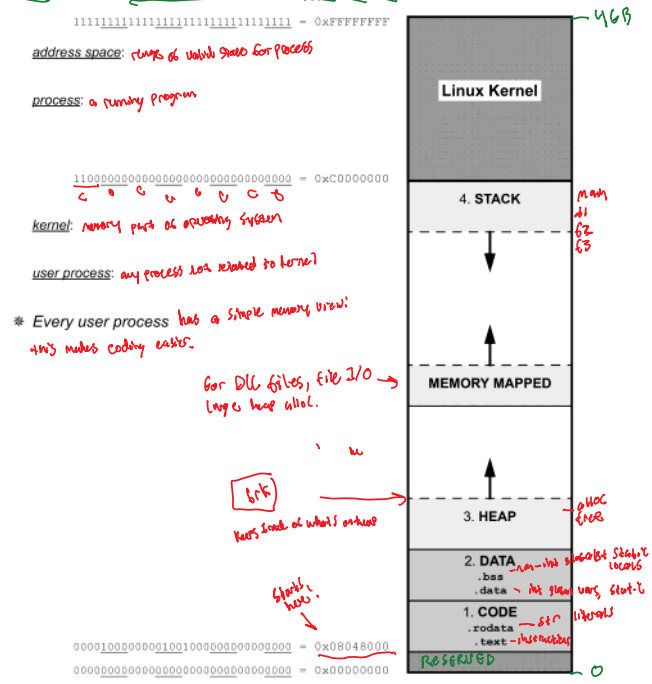

## C's Abstract Memory Model

1. **CODE Segment**
   - Contains: the program
     - **.text** section → binary machine code of program
     - **.rodata** section string literals go here
   - Lifetime: entire program's execution
   - Initialization: from executable file, done by loader before execution begins
     ↳ part of OS that knows how to start executable
   - Access: read-only

2. **DATA Segment**
   - Contains: Global variables, static local vars
   - Lifetime: entire program's execution
   - Initialization: from executable file, done by LOADER before exec begins.
     - **.data** section vars with non-zero init values
     - **.bss** section vars with 0 / non-init vars
       init
   - Access: read/write

3. **HEAP** (AKA Free Store)
   - Contains: memory that is ALLOC'D and FREE'D by program during exec
   - Lifetime: up to programmer (malloc/calloc/realloc) until it is freed
   - Initialization: N/A by default
   - Access: read/write

4. **STACK** (AKA Auto Store)
   - Contains: organized in stack frames → automatically alloc'd and free'd as functions execute
     *stack frame* (AKA activation record)
     non-static local vars, params, temp vars
   - Lifetime: scope (declaration → end of scope)
   - Initialization: NONE by default
   - Access: read/write

## Meet Globals and Static Locals

**What?**

A *global variable* is
- decl. outside any function
- access to all functions in file
- allocated in DATA segment

A *static local variable* is
- decl. inside a fcn with modifier "static"
- accessible only in its fcn
- allocated in Data segment before execution begins

**Why?**

for storage that exists during entire program execution

❄ *In general, global variables* SHOULD NOT
*Instead use* local vars that pass returned values to CALLER fncs

**How?**

```
#include <stdio.h>
int g = 11;                    // global

void f1(int p) {
  static int x = 22;           // static local, init once before execution begins
  x = x + p * g;
  printf("%d\n", x);
}

int main(void) {
  f1(g);
  g = 2;
  int g = 1;                   // local, shadows global of the same name
  f1(g);                       // g refers to local
  return 0;
}
```

*shadowing*: when local has same name as global
both access global var

❄ *Avoid shadowing; don't use the same identifier*

## Where do I live?

→ Identify the segment (and section) for each memory allocation in the code below.

```
#include <stdio.h>
#include <stdlib.h>

int gus = 14;
int guy;

int madison(int pam) {

    static int max = 0;
    int meg[] = {22,44,88};
    int *mel = &pam;
    max = gus--;
    return max + meg[1] + *mel;
}

int *austin(int *pat){

    static int amy = 33;
    int *ari = malloc(sizeof(int)*44);
    gus--;
    *ari = *pat;
    return ari;
}

int main(int argc, char *argv[]) {

   int vic[] = {33,66,99};
   int *wes = malloc(sizeof(int));
   *wes = 55;
   guy = 66;
   free(wes);
   wes = vic;
   wes[1] = madison(guy);
   wes = austin(&gus);
   free(wes);
   printf("Where do I live?");
   return 0;
}
```

- *Arrays, structs, and variables*

## Linux: Processes and Address Spaces

### Process and Job Control

- Linux is a multi-tasking OS where you can run jobs concurrently

`ps` lists snapshot of user processes

`jobs` lists only processes started from this line

`&` puts process in background

ctrl+z suspend proc

`bg` put suspended proc in background

`fg` put in foreground

ctrl+c break

`top` dynamic table of process resource usages

### Program Size

size <executable or object_file> display size of prog's mem segments + total size

a.out → compiled size (not running)

```
$gcc -m32 myProg.c
$size a.out
```

TOTAL

| text | data | bss | dec | hex | filename |
|---|---|---|---|---|---|
| 1029 | 276 | 4 | 1309 | 51d | a.out |

code data

### Virtual Address Space Maps

- Linux enables you to see the VAS (mem map) of each process

```
$pmap <pid_of_process>

$cat /proc/<pid_of_process>/maps

$cat /proc/self/maps
```

/proc:





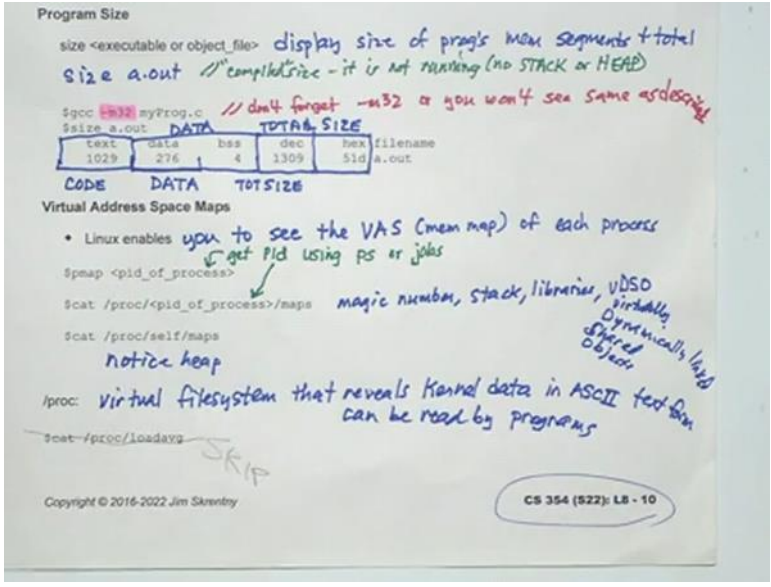

| | | |
|---|---|---|
| [ ] Lec 001 9:30am TR<br>[ ] Lec 002 1:00pm TR<br>Lecture | Print NetID | Print Name (first last) |

Computer Sciences 354

# Midterm Exam 1 Secondary

Thursday, October 6th, 2022

60 points (15% of final grade)

Instructors: Debra Deppeler

1. MARK an X in box by your lecture number above.
2. PRINT your NETID (UW login name not your photo id number) in box above.
3. PRINT your first and last name in box above.
4. FILL-IN all fields and their bubbles on the scantron form (use # 2 pencil).
   (a) LAST NAME - fill in your last (family) name starting at leftmost column.
   (b) FIRST NAME - fill in first five letters of your first (given) name.
   (c) IDENTIFICATION NUMBER is your UW Student ID number.
   (d) Under ABC of SPECIAL CODES, write your lecture number as a three digit value 001 or 002.
   (e) Under F of SPECIAL CODES, write the number 2 for Secondary and fill in the number (2) bubble.
5. DOUBLE-CHECK THAT YOU HAVE FILLED IN ALL ID FIELDS and that you have FILLED IN ALL CORRESPONDING BUBBLES ON SCANTRON.
6. Taking this exam indicates that you agree: to not write answers in large letters and to keep your answers covered; to not view or use another's work or any unauthorized devices in any way; to not make any type of copy of any portion of this exam; and that you understand that being caught doing any of these actions, or other actions that permit any student to submit work that is not wholly their own will result in automatic failure of the exam and possible failure of the course. Penalties are reported to the Deans Office for all involved.

| Parts | Number of Questions | Question Format | Possible Points |
|---|---|---|---|
| I | 10 | Single Choice | 20 |
| II | 10 | Multiple Choice | 30 |
| III | 2 | Written | 10 |
| | 22 | Total | 60 |

Assumptions unless instructions explicitly state otherwise:
addresses and integers are 4 bytes.
code questions are about C and IA-32/x86 assembly code on our Linux platform.

## Reference: Powers of 2

$2^5 = 32, 2^6 = 64, 2^7 = 128, 2^8 = 256, 2^9 = 512, 2^{10} = 1024$
$2^{10} = K, 2^{20} = M, 2^{30} = G$
$2^A * 2^B = 2^{A+B}, 2^A / 2^B = 2^{A-B}$

Turn over and put away all electronic devices and wait for the proctor to signal the start of the exam.

 CS354 Exam 1S 

 