

outlineL22-w11TR-student

Thursday, November 17, 2022 12:57 PM



outlineL22-
w11TR-st...

CS 354 - Machine Organization & Programming Tuesday Nov15, Thursday Nov 17, 2022

Exam Results expected by Friday Nov 18

Homework hw6: DUE on or before Monday Nov 21

Homework hw7: DUE on or before Monday Nov 28

Project p5: DUE on or before Friday Dec 2

Last Week

Instructions - SET Instructions - Jumps Encoding Targets Converting Loops	The Stack from a Programmer's Perspective The Stack and Stack Frames Instructions - Transferring Control Register Usage Conventions Function Call-Return Example
--	--

This Week

Function Call-Return Example (L20 p7) Recursion Stack Allocated Arrays in C Stack Allocated Arrays in Assembly Stack Allocated Multidimensional Arrays	Stack Allocated Structs Alignment Alignment Practice Unions
Next Week: Pointers in Assembly, Stack Smashing, and Exceptions B&O 3.10 Putting it Together: Understanding Pointers 3.12 Out-of-Bounds Memory References and Buffer Overflow 8.1 Exceptions 8.2 Processes 8.3 System Call Error Handling 8.4 Process Control through p719	

Recursion

Use a stack trace to determine the result of the call **fact (3)** :

```
int fact(int n) {
    int result;
    if (n <= 1) result = 1;
    else      result = n * fact(n - 1);
    return result;
}
```

direct recursion

recursive case

base case

"infinite" recursion

Assembly Trace

```
fact:
    pushl %ebp
    movl %esp, %ebp
    pushl %ebx
    subl $4, %esp


    movl 8(%ebp), %ebx
    movl $1, %eax

    cmpl $1, %ebx
    jle .L1

    leal -1(%ebx), %eax
    movl %eax, (%esp)
    call fact

    imull %ebx, %eax

.L1:
    addl $4, %esp
    popl %ebx
    popl %ebp
    ret
```



* *"Infinite" recursion causes*

* *When tracing functions in assembly code*

Stack Allocated Arrays in C

Recall Array Basics

$T \ A[N];$ where T is the element datatype of size L bytes
and N is the number of elements



1. Contiguous region of stack $L * N$ bytes
2. Identifier associated with the start of the array

* *The elements of A* Are accessed using address arithmetic

Recall Array Indexing and Address Arithmetic

L = element size

$X(a)$ = start address of the array

$$\&A[i] \equiv A + i = X(a) + L * i$$

(index)

→ For each array declarations below, what is L (element size), the address arithmetic for the i th element, and the total size of the array?

C code	L	address of i th element	total array size
1. <code>int I[11]</code>	4	$X(i) + 4i$	44
2. <code>char C[7]</code>			
3. <code>double D[11]</code>			
4. <code>short S[42]</code>	2	$X(s) + 2i$	84
5. <code>char *C[13]</code>			
6. <code>int **I[11]</code>	4	$X(i) + 4i$	44
7. <code>double *D[7]</code>			

Stack Allocated Arrays in Assembly

Arrays on the Stack

→ How is an array laid out on the stack? Option 1 or 2:

* *The first element (index 0) of an array*

Is closest to the top of the stack.

Accessing 1D Arrays in Assembly


IA-32 is designed to simplify array access

Assume array's start address in `%edx` and index is in `%ecx`

`movl (%edx, %ecx, 4), %eax` \equiv $M[x(a) + 4*i] = A[i]$
 In C: `*(A+i)`

→ Assume `I` is an `int` array, `S` is a `short int` array, for both the array's start address is in `%edx`, and the index `i` is in `%ecx`. Determine the element type and instruction for each:

C code	type	assembly instruction to move C code's value into <code>%eax</code>
1. <code>I</code>	<code>int*</code>	<code>X(I)</code> <code>Movl %edx, %eax</code>
2. <code>I[0]</code>	<code>int</code>	<code>M[X_i] Movl (%edx, %ecx), %eax</code>
3. <code>*I</code>		
4. <code>I[i]</code>		
5. <code>&I[2]</code>	<code>int*</code>	<code>leal 8(%edx, %ecx), %eax</code>
6. <code>I+i-1</code>	<code>int*</code>	<code>leal -1(%edx, %ecx), %eax</code>
7. <code>*(I+i-3)</code>		
8. <code>S[3]</code>		
9. <code>S+1</code>		
10. <code>&S[i]</code>		
11. <code>S[4*i+1]</code>	<code>short*</code>	<code>movsl 4(%edx, %ecx, 8), %eax</code>
12. <code>S+i-5</code>		




Stack grows always to the lower address

Stack Allocated Multidimensional Arrays

Recall 2D Array Basics

T A[R][C]; where T is the element datatype of size L bytes.
 R is the number of rows and C is the number of columns



* Recall that 2D arrays are stored on the stack in row major order

```
int A[5][3];          typedef int row_t[3];
                      row_t A[5];
```

col per row

rows

Accessing 2D Arrays in Assembly

$\&A[i][j] \equiv$ Start address + offset to row i + offset to col j .

$$X_A + (C * i * L) + (L * j)$$

Given array A as declared above, if X_A in $\%eax$, i in $\%ecx$, j in $\%edx$ then $A[i][j]$ in assembly is:

```
leal (%ecx, %ecx, 2), %ecx    3i in %ecx
sall $2, %edx                4j
addl %eax, %edx              (X_A + 4j)
movl (%edx, %ecx, *4), %eax   X_A + (3i*4) + 4j -> %eax
```

Compiler Optimizations

- ♦ If only accessing part of array

Compiler points to that part of the array

Then offset from there

- ♦ If taking a fixed stride through the array

Compiler uses stride * element size as offset

Stack Allocated Structures

Structures on the Stack

```

struct iCell {
    int x;
    int y;
    int c[3];
    int *v;
};

```

Handwritten notes for the struct definition:


# byte	offset	Hex
4	0	0x0
4	4	0x4
12	8	0x8
4	20	0x14

→ How is a structure laid out on the stack? Option 1 or 2:

The compiler

- Associates data member names with the Offset of the Struct, from the start of the struct

- Use address arithmetic with offset to access data members



* The first data member of a structure is closest to the top of the stack

Accessing Structures in Assembly

Given:

```

→ struct iCell ic = //assume ic is initialized
void function(iCell *ip) {

```

→ Assume ic is at the top of the stack, %edx stores ip and %esi stores i. Determine for each the assembly instruction to move the C code's value into %eax:

C code	assembly
1. ic.v	movl 20(%esp), %eax
2. ic.c[i]	
3. ip->x	movw (%edx), %eax
4. ip->y	
5. &ip->c[i]	leal 8(%edx, %esi, 4), %eax


* Assembly code to access a structure does not have data member names, only Offsets!

Alignment

What? Most systems restrict the addresses where primitive data can be stored.

Why? Primarily for performance + cheaper hardware

Example: Assume cpu reads 8 byte words
f is a misaligned float



It's slow: misalignment can take multiple reads
2 reads, extracts, combines

Restrictions

IA-32 has no alignment restrictions

Linux: short Must be multiple of two
int, float, pointer, double Must be multiple of four

Windows: same as Linux except
double Must be multiple of eight


Implications

Going to add padding to make the alignment correct by the compiler

Padding vs. no Padding

Structure Example

```
struct s1 {  
    int i;  
    char c;  
    int j;  
};
```



* **The total size of a structure** is typically a multiple of its largest data member size
-In this diagram it would be four, so it's a multiple of four

Alignment Practice

→ For each structure below, complete the memory layout and determine the total bytes allocated.

```
1) struct sA {  
    int i;  
    int j;  
    char c;  
};
```



```
2) struct sB {  
    char a;  
    char b;  
    char c;  
};
```



```
3) struct sC {  
    char c;  
    short s;  
    int i;  
    char d;  
};
```



```
4) struct sD {  
    short s;  
    int i;  
    char c;  
};
```



```
5) struct sE {  
    int i;  
    short s;  
    char c;  
};
```



* *The order that a structure's data members are listed*
Can affect memory utilization

Copyright © 2016-2022 Jim Skrentny

CS 354 (F22): L22 - 8

Unions

What? A union is

- ♦ Like a struct except fields share same memory
- ♦ Allocates only enough memory for largest field


Why?

- ♦ Allows data to be accessed as different types
- ♦ Used to access hardware
- ♦

How? Implement "polymorphism"


How?

How? Implement "polymorphism"



Example

```
typedef union {
  unsigned char cntrlrByte;
  struct {
    unsigned char playbutn : 1;
    unsigned char pausebutn : 1;
    unsigned char ctrlbutn : 1;
    unsigned char fire1butn : 1;
    unsigned char fire2butn : 1;
    unsigned char direction : 3;
  } bits;
} CntrlrReg;
```



CntrlrReg c1;

readctrlReg(c1.cntrlrByte); //this reads all 8 bits into c1, hypothetically

If (c1.bits.fire1butn) { ... fireshot(); } //we are using the union here to make it convenient for different uses